

Foundation models *for spatial-time series*

Week 2

Classic versus Vibe coding

- | | |
|------------------------------------|---------------------------------------|
| 1. Your project to develop | 1. Your project among alternatives |
| 2. Computational experiment setup | 2. Prompt engineering |
| 3. Data download and preprocessing | 3. Library datasets |
| 4. Your model | 4. Your <i>generated</i> model |
| 5. Alternative models | 5. Generated superposition of models |
| 6. Computational experiment | 6. Generated model comparison |
| 7. Manual reporting | 7. Generated report and paper |

The Vibe Coding Technique

1. Start with curiosity, not software architecture

Instead of planning a whole system, just pick a question

The vibe is to play, not to ship production code

2. Keep it lightweight

A Jupyter notebook, Google Colab, or a single Python file

Only import the minimum you need

3. Prototype in the smallest way possible

Instead of building a whole dataset loader,
hardcode a tiny dataset

4. Write tiny models first

Instead of aiming for GPT-4, write a baby transformer

5. Iterate by feel

Once something works, ask “what if I add...?”

Add self-attention. Add positional encodings.

Scale up dataset a bit.

Each step should be incremental and driven by curiosity.

6. Embrace imperfection

Code will be messy. That’s okay — vibe coding is about insight first, refactor later.

When the experiment feels solid, *then* clean it up.



Motivation

As we write papers,

we make a lots of code to compare many similar models.

The Vibe coding style

eases the process of model collection, creation, and comparison.

Minimum problem statement

There given as input

- 1) a spatial time series and its context (physics of measurements)
- 2) a mathematical, computational or machine learning model (prior)
- 3) a prompt

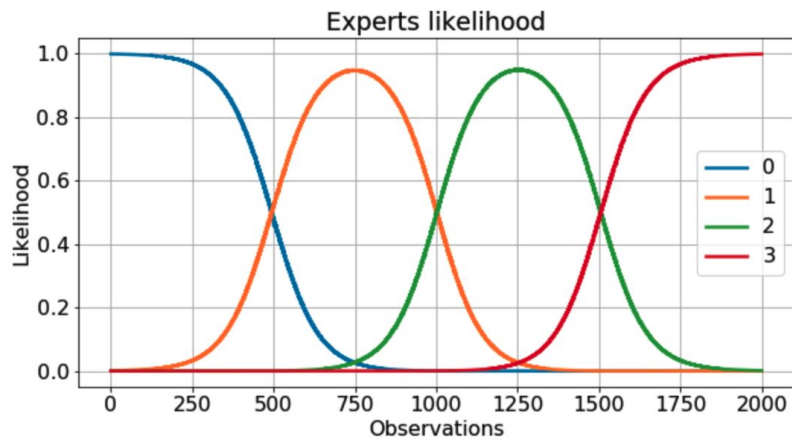
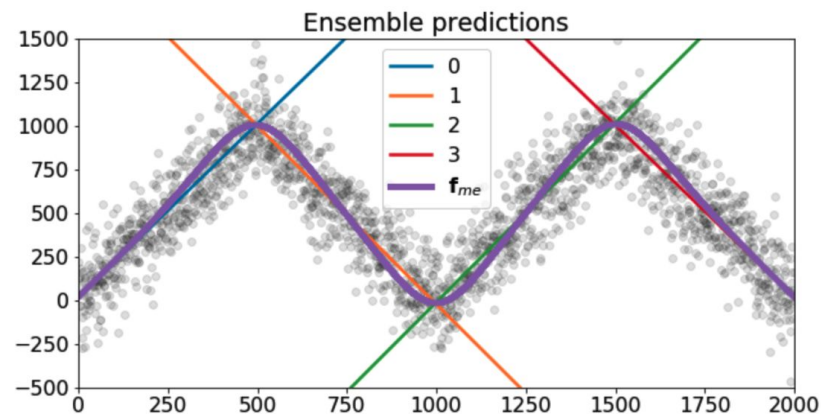
Expected as output

- 1) a selected model that fits data
- 2) an analytical report to the prompt


Prompt examples

1. Find change points in the [time series]
2. Select a model from [class] for [time series] given [mathematical model]

PINN is the Hinton's distillation of models



Creating composite loss function



$$f = \mathbf{w}_{1 \times 1_k}^T \mathbf{z}_{k-1} \circ \mathbf{W}_{k-1}^T \mathbf{z}_{k-2} \circ \cdots \circ \mathbf{W}_2^T \mathbf{z}_1 \circ \mathbf{W}_1^T \mathbf{x}$$

$n_2 \times 1$ $n_1 \times n$ $n \times 1$

Neural network error

$$E_y = (y_i - f(\mathbf{x}))^2$$

Autoencoder reconstruction error

$$E_x = \|\mathbf{x} - \mathbf{r}(\mathbf{z})\|^2$$

Types of autoencoders

PCA

$$\mathbf{W}^T \mathbf{W} = \mathbf{I}_n$$

skip block

$$\mathbf{W} = \mathbf{I}_n$$

metric

$$\mathbf{x}^T \mathbf{W} \mathbf{x} \geq 0$$

multi-linear

$$\underline{\mathbf{W}} \mathbf{X}$$

Autoencoder transform: $\mathbf{z} = (1 + \exp(-\mathbf{W}^T \mathbf{x} + \mathbf{b}))^{-1}$

Reconstruction decoder: $\hat{\mathbf{x}} = \mathbf{r}(\mathbf{z}(\mathbf{x}))$

loss → model → parameters → optimizer

To construct a tool, one has to generate the composite loss function

```
model = TinyGPT(vocab_size)
optimizer = torch.optim.AdamW(model.parameters(), lr=1e-3)

for step in range(200): # try 2000+ for better results
    xb, yb = get_batch()
    logits, loss = model(xb, yb)
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    if step % 50 == 0:
        print(step, loss.item())
```


Knowledge distillation between physics model and NN

For some objects \mathbf{x} there given **privileged** information \mathbf{x}^* . Introduce a **student** model $\mathbf{f}_s \in \mathfrak{F}_s$ and a **teacher** model $\mathbf{f}_t \in \mathfrak{F}_t$:


$$\mathbf{f}_s : \mathbf{x} \longrightarrow \mathbf{y}, \quad \mathbf{f}_t : \mathbf{x}^* \longrightarrow \mathbf{y}.$$

$$\mathbf{f}_s = \arg \min_{\mathbf{f} \in \mathfrak{F}_s} \frac{1}{n} \sum_{i=1}^n \left[(1 - \lambda) S(\mathbf{y}_i, \mathbf{f}(\mathbf{x}_i)) + \lambda S(\mathbf{s}_i, \mathbf{f}(\mathbf{x}_i)) \right],$$

T is the temperature, parameter of smoothing, for classification:


$$\mathbf{s}_i = \boldsymbol{\sigma}(\mathbf{f}_t(\mathbf{x}_i)/T), \quad S(\mathbf{y}_i, \mathbf{f}(\mathbf{x}_i)) = - \sum_{k=1}^c \mathbf{y}_k \log \boldsymbol{\sigma}(\mathbf{f}(\mathbf{x}_i)), \quad \boldsymbol{\sigma} - \text{softmax}$$

Time series and the state space embedding



$t = 1000$

$\dim(s) \approx 1000$




$\dim(x) = 4$

Reduce dimensionality with the principal component analysis, autoencoder $z = \mathbf{W}^T \mathbf{x}$ where \mathbf{W} is an orthogonal (rotation) matrix. The first principal components are given by Singular Values Decomposition

$$\sqrt{\lambda_k} \mathbf{v}_k = \mathbf{X}^T \mathbf{u}_k \quad \text{the SVD is} \quad \mathbf{X} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$$

Forecasting problem and model selection



$x \rightarrow x$ transformations

- ① Linear model
- ② Stack of autoencoders
- ③ Neural ODE
- ④ Neural PDE
- ⑤ Graph diffusion

$y \rightarrow y$ transformations

- ① Linear model
- ② Spherical harmonics
- ③ Topological alignment
- ④ Lagrangian, Hamiltonian neural networks

Initial models

1. Models

Direct models: AR, ARIMA, GRU, LSTM

Metric models: LLE, DM, GH, RBF

Non-parametric: GPR

2. Ways to construct state spaces

SSM models: S4, S5, Hippo

Kalman

SSA, SSM

3. Ways to transform state spaces

FT, OL, ODE

CCA, CCM

The time series has two domains: the time domain and the frequency domain. The spatial time series also has a metric space or metric tensor that changes in time.

Or any model from Prophet: <https://github.com/facebook/prophet>

For the Week 3

1. Join a group
2. Discuss the goals of the project and a solution (see the problem statement above)
3. Make a review of various ways to solve the problem
4. Select an LLM, GPT
5. Run the code to check if it works
 1. Store the code in the group repository
 2. Store the talk slides/report, too
6. Make a 10-minute talk about
 1. Functionality and architecture/principles of the model
 2. Why did you select this model
 3. The alternative models to select from

The challenge

How to generate the composite loss function?

How to manage hyperparameters (components of the loss function)?

Do we need to treat the tokens (what LLM forecasts) as hyperparameters?

How to select a class of models according to request?

How to select a model from a class?

`__init__()`

Do we need to use tool-embedding class of LLM or there are the other solutions?

To read

Papers

1. *Konstantin Yakovlev et al. Toolken+: Improving LLM Tool Usage with Reranking and a Reject Option, 2025*
2. *Shibo Hao et al. ToolkenGPT: Augmenting Frozen Language Models with Massive Tools via Tool Embeddings, 2023*
3. *Hugo Touvron et al. LLaMA: Open and Efficient Foundation Language Models, 2023*

Sandbox examples

1. [Token_Example.ipynb](#)
2. [TinyGPT_from_Scratch.ipynb](#)
3. <https://github.com/karpathy/nanoGPT>
4. <https://github.com/facebook/prophet>

- EXPLORE
- LLMS-FROM-SCRATCH-MAIN
 - .github
 - ISSUE_TEMPLATE
 - workflows
 - appendix-A
 - appendix-D
 - appendix-E
 - ch01
 - ch02
 - 01_main-chapter-code
 - .ipynb_checkpoints
 - ch02-checkpoint.ipynb
 - dataloader-checkpoint.ipynb
 - exercise-solutions-checkpoint.ipynb
 - README-checkpoint.md
 - the-verdict-checkpoint.txt
 - ch02.ipynb
 - dataloader.ipynb
 - exercise-solutions.ipynb
 - README.md
 - the-verdict.txt
 - 02_bonus_bytepair-encoder
 - gpt2_model
 - encoder.json
 - vocab.bpe
 - bpe_openai_gpt2.py
 - compare-bpe-tiktoken.ipynb
 - README.md
 - requirements-extra.txt
 - 03_bonus_embedding-vs-matmul
 - 04_bonus_dataloader-intuition
 - 05_bpe-from-scratch
 - README.md

ch02 > 01_main-chapter-code > ch02.ipynb > M+ Chapter 2: Working with Text Data > M+ empty cell

Generate + Code + Markdown | Run All Clear All Outputs | Outline ...

Supplementary code for the [Build a Large Language Model From Scratch](#) book by [Sebastian Raschka](#)

Code repository: <https://github.com/rasbt/LLMs-from-scratch>

Chapter 2: Working with Text Data

Packages that are being used in this notebook:

```
from importlib.metadata import version

print("torch version:", version("torch"))
print("tiktoken version:", version("tiktoken"))
```

[1]

... torch version: 2.5.1
tiktoken version: 0.7.0

Generate + Code

- This chapter covers data preparation and sampling to get input data "ready" for the LLM

